

Fully functioning Web application using serverless technologies

AWS Lambda

API Gateway

Amazon S3

Amazon DynamoDB

Amazon Simple Email Service (SES)

IAM Roles

1. Project Overview

Company Name: Belly Brew Co. Limited

Industry: Manufacturing and distribution of gut-friendly beverages and ready meals

Market: Retail supermarkets, grocery stores, restaurants, and coffee shops across the United Kingdom

Mission Statement: Innovating gut-friendly meals and beverages that fuse nutrition and convenience, fostering well-being with every choice.

2. Problem Statement

Belly Brew Co. Limited seeks to expand customer awareness and brand loyalty while capturing additional market share in the ready meals and soft drinks industry. The company aims to engage customers by providing an enticing offer through a website, encouraging users to sign up for a newsletter in exchange for a free gut-friendly recipe eBook. The ultimate goal is to create a secure, scalable, and manageable cloud-based solution to gather customer details and facilitate email marketing campaigns.

3. Business Requirements

Primary Goal:

- **Customer Engagement and Market Expansion:** Capture the names and email addresses of both existing and new customers interested in Belly Brew Co.'s product range.
- **Brand Awareness:** Build brand loyalty and customer awareness across the ready meals and beverages industry.

Key Requirements:

1. **Customer Information Capture:** A web-based solution to collect names and email addresses from customers.
 2. **Newsletter Subscription Form:** The website should have a subscription form that entices users to sign up for the newsletter in exchange for a free gut-friendly recipe eBook.
 3. **Email Campaigns:** The gathered email addresses will be used for future marketing campaigns.
 4. **E-book Delivery:** Once the user signs up for the newsletter, the system must deliver a free recipe eBook via email.
-

4. Key Stakeholder Insights

- **Ellie Brew (CEO):**
The company's marketing department requires a cloud solution to capture customer details (name and email). The solution should work in collaboration with Belly Brew Co.'s internal development team.
- **Anisha Belly (Head of Marketing):**
The marketing team wants a website with a subscription form for visitors to sign up for the newsletter and receive a free eBook with over 100 recipes. The captured email addresses will be used for future marketing campaigns.
- **George Gut (CTO):**
The company seeks a cloud-based, serverless solution that is secure, scalable, and requires no administrative or management overhead. The solution should ensure that only genuine signups

receive the eBook, avoiding fraudulent or automated submissions. The company also aims to avoid the overhead of managing and maintaining servers.

5. Technical Requirements

Solution Characteristics:

- **Serverless Architecture:** A serverless solution will avoid the administrative overhead of managing servers and ensure automatic scaling, security, and resilience.
 - **Data Storage:** Customer details (names and email addresses) should be securely stored in a database for marketing purposes.
 - **Email Automation:** A system must be in place to automatically email the eBook to subscribers immediately after sign-up.
 - **Security:** The system must validate that only genuine users who sign up receive the eBook.
-

6. Solution Design Overview

High-Level Workflow:

1. **User Interaction:** Customers, such as Gemma, visit the Belly Brew Co. website and subscribe to the newsletter by submitting their name and email.
 2. **Form Submission:** Upon submission, the data is processed by a serverless compute service that stores customer details in a database (e.g., DynamoDB) for marketing purposes.
 3. **Email Delivery:** An automated email is sent to the user with a secure link to download the recipe eBook.
 4. **E-Book Access:** The customer receives the eBook and can download it from the link provided in the email, ensuring a positive user experience while securing the file from unauthorized access.
-

7. Conclusion

The proposed solution for Belly Brew Co. Limited addresses their need to capture customer data and deliver a value proposition through an automated, serverless cloud architecture. By implementing this solution, Belly Brew Co. can increase customer engagement, market share, and brand loyalty while keeping operational costs low and security high.

1. Reducing Infrastructure Management Overhead

- **Serverless** technologies like AWS Lambda are utilized to reduce the need for managing infrastructure.
- **Amazon S3** is used to host the static website and the eBook itself, avoiding the need for EC2 instances or managing servers.
- **Amazon DynamoDB** provides a no-SQL database solution, avoiding the need for database administration (such as with RDS).
- **Amazon API Gateway** facilitates communication between frontend and backend services, decoupling components and ensuring scalable, manageable APIs.

2. Ensuring Secure eBook Delivery

- The eBook is stored in an **S3 bucket** and accessed via a **pre-signed URL** generated dynamically.
- This URL is time-limited, preventing widespread sharing while still providing user-friendly access.
- **Amazon Simple Email Service (SES)** handles the email functionality, securely sending the pre-signed URL link to users.

3. Detailed Process Overview

- **User Interaction:** The customer submits a form via the static website.
- **Lambda Functions:** Triggered by API Gateway, Lambda performs business logic, stores user details in DynamoDB, and generates the pre-signed URL.

- **Email Delivery:** SES sends the pre-signed URL, ensuring a smooth flow from signup to delivery with time-restricted access to the eBook.

4. Additional Considerations

- **Security:** IAM roles are configured to ensure each service (Lambda, API Gateway, S3, etc.) has the necessary permissions, keeping the application secure.
- **Cost Efficiency:** Using serverless technologies like Lambda and S3 hosting allows for cost savings, especially when services are idle.

This architecture ensures that you meet the key requirement of reducing infrastructure overhead while still delivering a secure, scalable solution for the eBook distribution. It also provides a clear separation of concerns between the frontend and backend, leveraging AWS's fully managed services to handle all key tasks without manual intervention.

To summarize the steps for setting up this serverless web application:

1. **Create an S3 bucket** for storing your recipe book. Make sure the bucket is private and upload your book.pdf file in the correct folder structure.
2. **Set up a DynamoDB table** for storing the user profiles (name and email). Use email as the partition key and leave the default capacity settings.
3. **Configure Amazon SES (Simple Email Service)** in a sandbox environment, and verify an email identity that will be used to send emails.
4. **Create an IAM Policy and Role** for AWS Lambda. The policy should allow actions on the S3 bucket, DynamoDB, and SES services, and should be attached to the Lambda role.
5. **Create a Lambda function** using Python 3.9. Paste the provided Lambda code and configure the necessary environment variables such as the email address, S3 bucket name, DynamoDB table name, and region.
6. **Set up API Gateway** with an HTTP API, integrating it with the Lambda function. Define the /submit path for the POST method, which will trigger the Lambda function when the form is submitted.

7. HTML & API Integration:

You modified the HTML front-end, specifically line 4 and 5, to include the API Gateway invoke URL (with /submit appended). This URL triggers the backend flow when the form on the static website is submitted, sending the user's data to API Gateway.

8. S3 Setup:

You created an S3 bucket to host your static website. After setting the correct region, you ensured that public access was allowed by unchecking the block public access option and acknowledging that the bucket would be publicly accessible. Then, you uploaded all the static website files (HTML, CSS, JS, images, etc.).

9. Permissions & Policy:

Although public access was enabled, you still needed to apply an appropriate bucket policy. Using the AWS Policy Generator, you created an S3 bucket policy that allows public access to GET the objects (files) in your bucket. You included both the bucket ARN and its contents in the policy to make everything accessible.

10. Static Website Hosting:

You enabled static website hosting in the bucket properties, specifying the **index** and **error** pages. The index.html serves as the main page, while the error page is displayed if something goes wrong (e.g., file not found).

11. Testing:

After setting up the website and integrating the API, you tested the form submission process by filling in the form (name and email) and clicking the submit button. This action triggered the API Gateway, which invoked the Lambda function.

12. Lambda Execution:

The Lambda function processed the form submission by:

- Inserting the user's name and email into a DynamoDB table.
- Generating a pre-signed URL for the ebook stored in S3.
- Using Amazon SES (Simple Email Service) to send an email to the user with the download link for the ebook.

13. Monitoring & Verification:

You checked the **Lambda** function's monitoring tab to confirm that it was

invoked when the form was submitted. You also verified in **DynamoDB** that the user's details were stored successfully, ensuring that the system captured the user's information as expected.

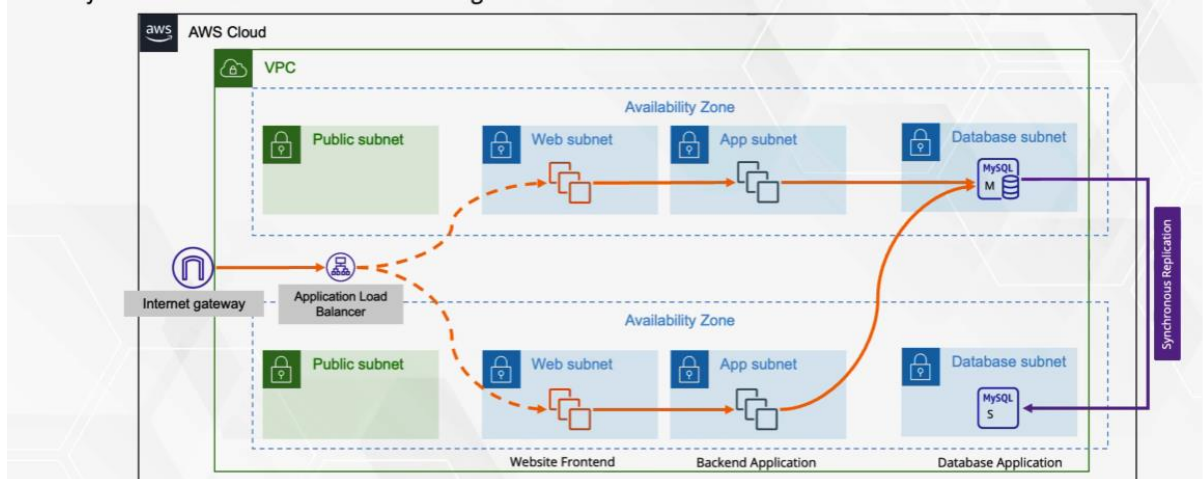
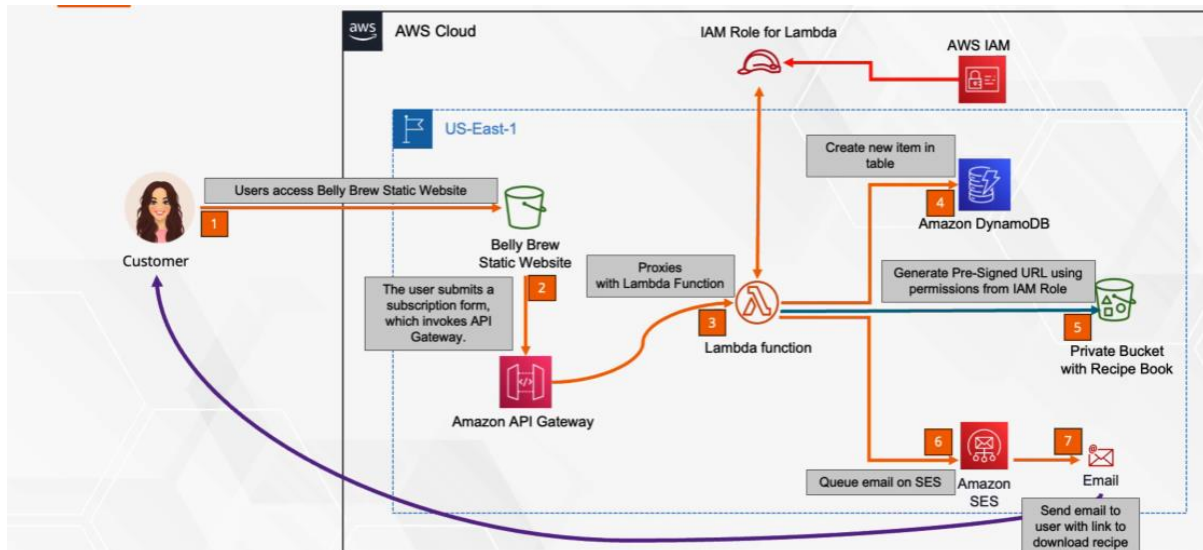
14. Final Output:

The user received an email containing a link to download the ebook, completing the process. This demonstrated how the entire serverless architecture—from front-end submission to backend processing—was working seamlessly.

Steps

1. Create Amazon S3 Bucket For Recipe Book
2. Create DynamoDB Table
3. Configure Simple Email Service (Amazon SES)
4. Define IAM Policy and Role for AWS Lambda
5. Create Lambda Function
6. Build HTTP API with API Gateway
7. Create Amazon S3 Static Website Hosting
8. Test Application

ARCHITECTURE



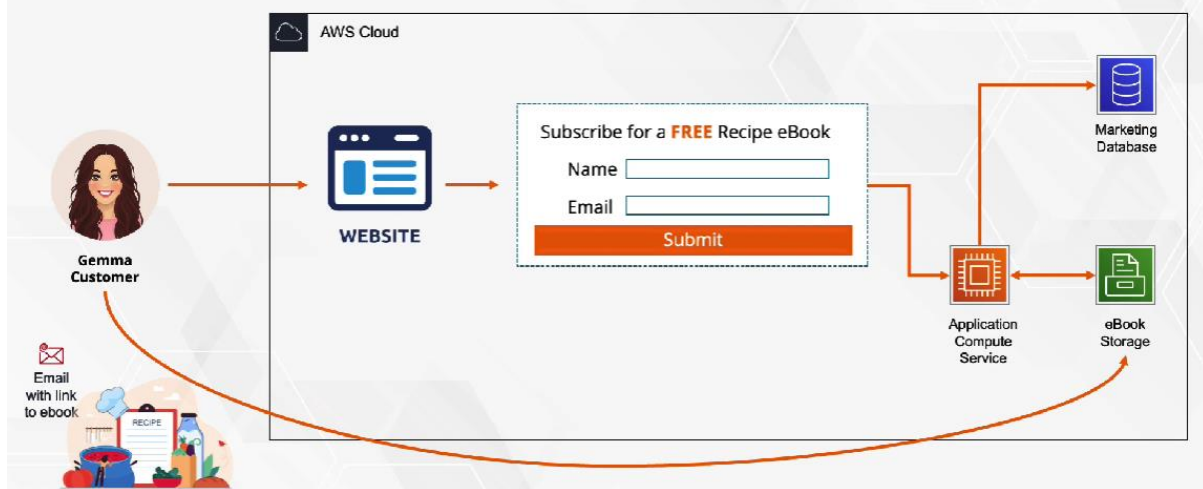
Introducing Amazon S3 Pre-Signed URLs

"secure, scalable and resilient, ensure that the eBook is only available to genuine sign-ups."

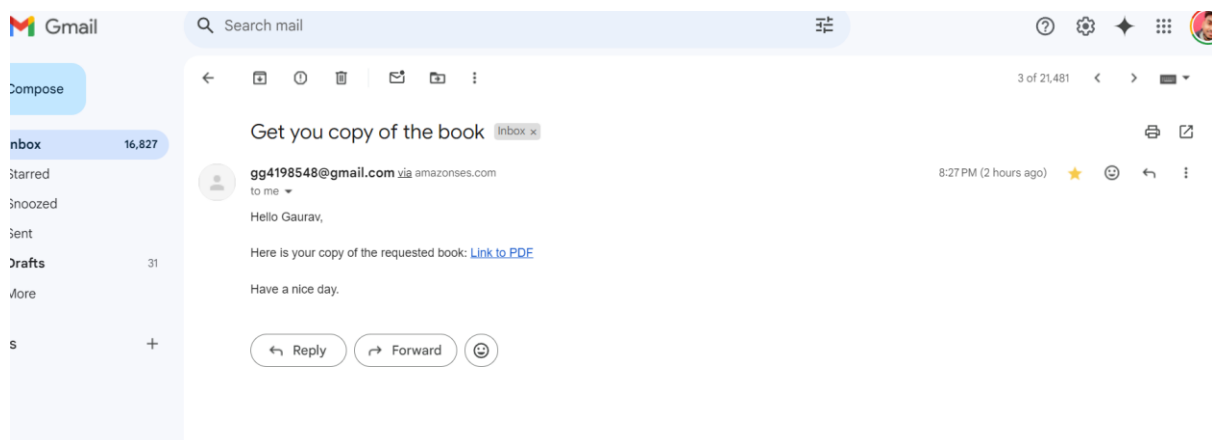
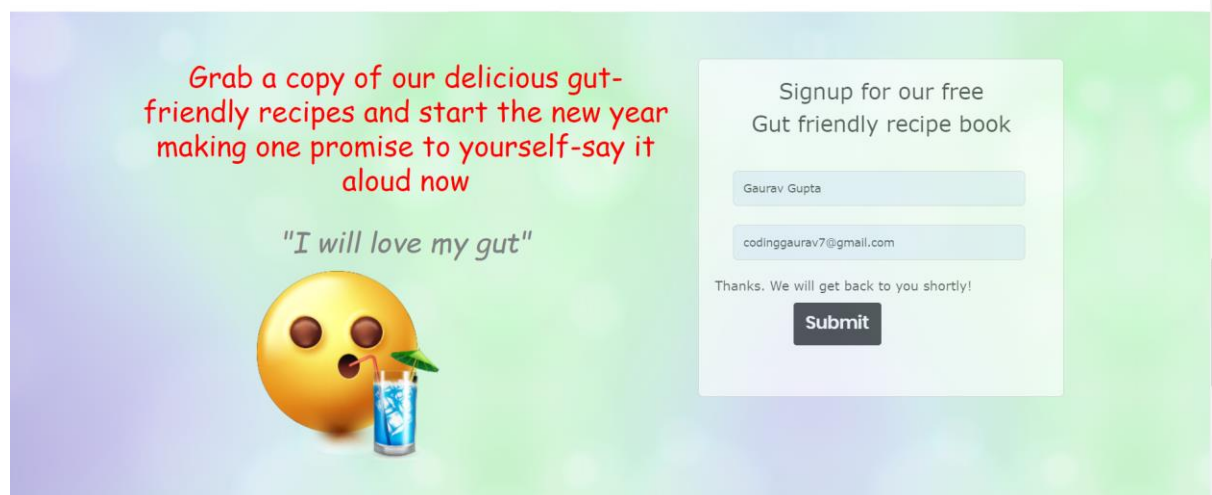



Fulfilling the business requirements

"secure, scalable and resilient, ensure that the eBook is only available to genuine sign-ups."



Deployed Website



S3

Amazon S3

Buckets

Access Grants

Access Points

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

Storage Lens groups

AWS Organizations settings

Amazon S3 > Buckets

Account snapshot - updated every 24 hours

All AWS Regions

View Storage Lens dashboard

General purpose buckets

Directory buckets

General purpose buckets (3)

Info

All AWS Regions

Copy

Copy ARN

Empty

Delete

Create bucket

Buckets are containers for data stored in S3.

Find buckets by name

< 1 > ⚙

	Name	AWS Region	IAM Access Analyzer	Creation date
<input type="radio"/>	bellybrewrecipebook-1	US East (N. Virginia) us-east-1	View analyzer for us-east-1	October 14, 2024, 19:32:28 (UTC+05:30)
<input type="radio"/>	bellywebapp	US East (N. Virginia) us-east-1	View analyzer for us-east-1	October 14, 2024, 20:18:32 (UTC+05:30)

amazon s3

Buckets

Access Grants

Access Points

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

Storage Lens groups

AWS Organizations settings

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (6)

Info

Refresh

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

< 1 > ⚙

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	css/	Folder	-	-	-
<input type="checkbox"/>	error.html	html	October 14, 2024, 20:22:04 (UTC+05:30)	10.0 B	Standard
<input type="checkbox"/>	fonts/	Folder	-	-	-
<input type="checkbox"/>	images/	Folder	-	-	-
<input type="checkbox"/>	index.html	html	October 14, 2024, 20:22:03 (UTC+05:30)	12.3 KB	Standard
<input type="checkbox"/>	js/	Folder	-	-	-

Static website hosting

Edit

Use this bucket to host a website or redirect requests. [Learn more](#)

S3 static website hosting

Enabled

Hosting type

Bucket hosting

Bucket website endpoint

When you configure your bucket as a static website, the website is available at the AWS Region-specific website endpoint of the bucket. [Learn more](#)

<http://bellywebapp.s3-website-us-east-1.amazonaws.com>



Step 1: Select Policy Type

A Policy is a container for permissions. The different types of policies you can create are an [IAM Policy](#), an [S3 Bucket Policy](#), an [SNS Topic Policy](#), a [VPC Endpoint Policy](#), and an [SQS Queue Policy](#).

Step 2: Add Statement(s)

Effect ☒ Allow ☐ Deny

Use a comma to separate multiple values.

AWS Service

☐ All Services (*)

Use multiple statements to add permissions for more than one service.

Actions -- Select Actions -- ☰ ☐ All Actions (*)

Amazon Resource Name (ARN)

Actions -- Select Actions -- ☐ All Actions ('*')

Amazon Resource Name (ARN)

ARN should follow the following format: `arn:aws:s3:::${BucketName}/${KeyName}`.
Use a comma to separate multiple values.

Add Conditions (Optional)

Add Statement

You added the following statements. Click the button below to Generate a policy.

Step 3: Generate Policy

A *policy* is a document (written in the [Access Policy Language](#)) that acts as a container for one or more statements.

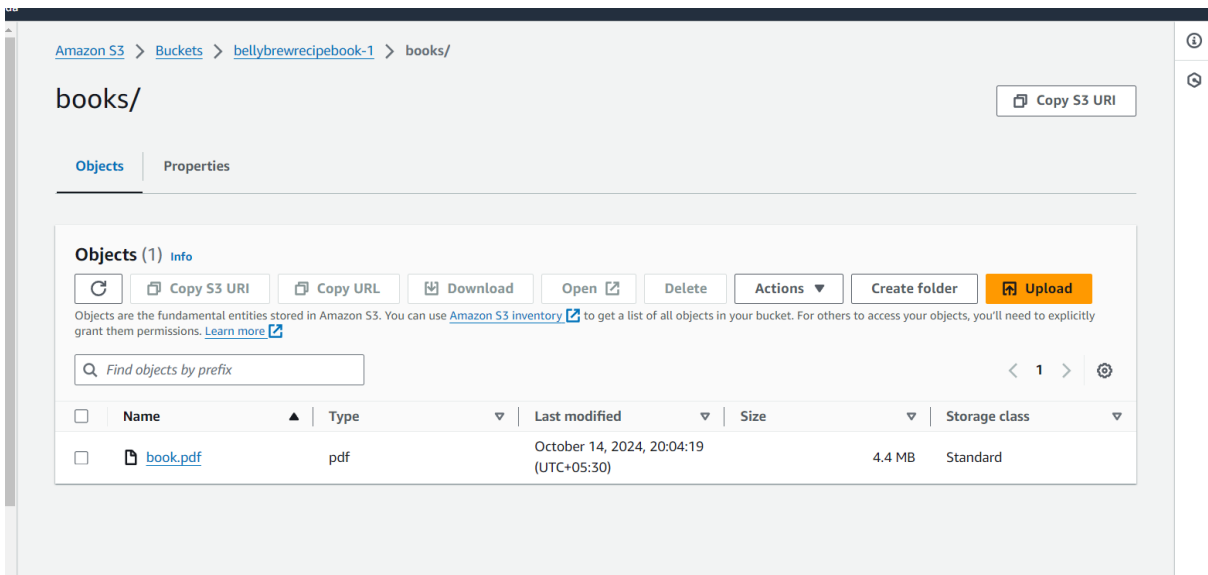
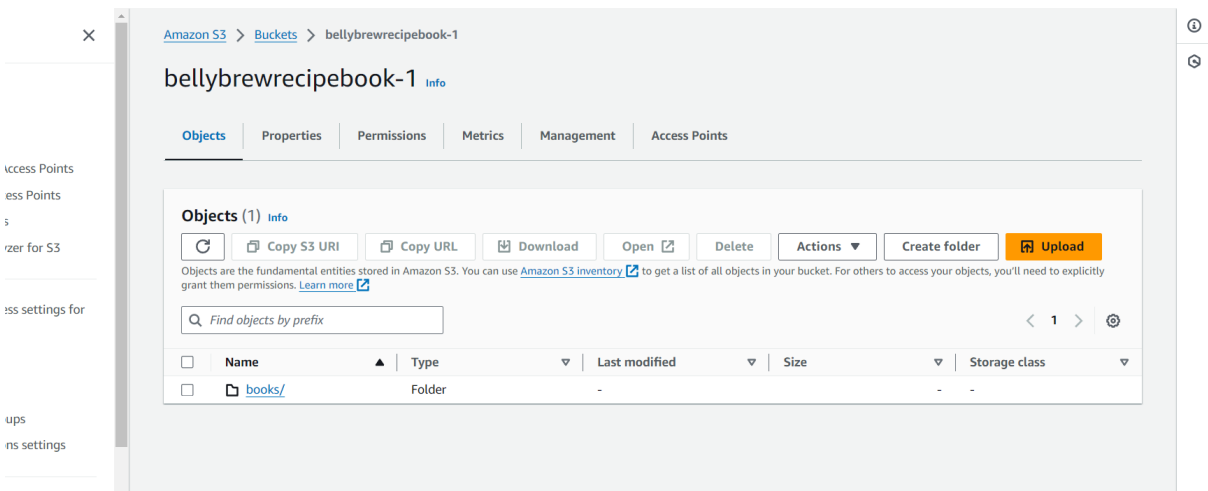
Generate Policy Start Over

This AWS Policy Generator is provided for informational purposes only, you are still responsible for your use of Amazon Web Services technologies and ensuring that your use is in compliance with all applicable terms and conditions. This AWS Policy Generator is provided **as is** without warranty of any kind, whether express, implied, or statutory. This AWS Policy Generator does not modify the applicable terms and conditions governing your use of Amazon Web Services technologies.

©2010, Amazon Web Services LLC or its affiliates. All rights reserved.

An **amazon.com** company

Books



IAM

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings

Access reports

- Access Analyzer
- External access
- Unused access
- Analyzer settings

Allows Lambda functions to call AWS services on your behalf.

Summary Edit

Creation date
October 14, 2024, 19:54 (UTC+05:30)

ARN
arn:aws:iam::569124580578:role/BellyBrewSubmitHandlerRole1

Last activity
2 hours ago

Maximum session duration
1 hour

Permissions | Trust relationships | Tags | Last Accessed | Revoke sessions

Permissions policies (1) Info Simulate Remove Add permissions

You can attach up to 10 managed policies.

Filter by Type
All types

Policy name	Type	Attached entities
BellyBrewSubmitHandlerPolicy1	Customer managed	1

Lambda

BellyBrewSubmitHandlerFunction-1

Throttle Copy ARN Actions

Function overview Info Export to Application Composer Download

Diagram Template

BellyBrewSubmitHandler Function-1

Layers (0)

API Gateway

+ Add trigger + Add destination

Description
-

Last modified
3 hours ago

Function ARN
arn:aws:lambda:us-east-1:569124580578:function:BellyBrewSubmitHandlerFunction-1

Function URL Info

Code | Test | Monitor | Configuration | Aliases | Versions

Dashboard

Applications

Functions

- BellyBrewSubmitHandlerFunction-1

Additional resources

- Code signing configurations
- Event source mappings
- Layers
- Replicas

Related AWS resources

- Step Functions state machines

File Edit Find View Go Tools Window Test Deploy

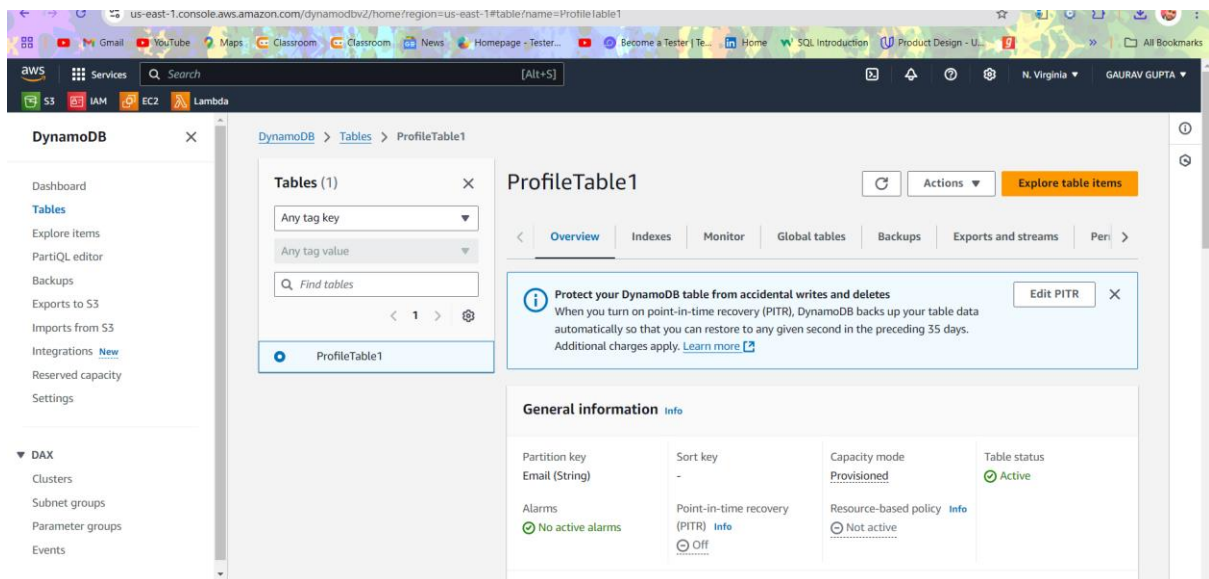
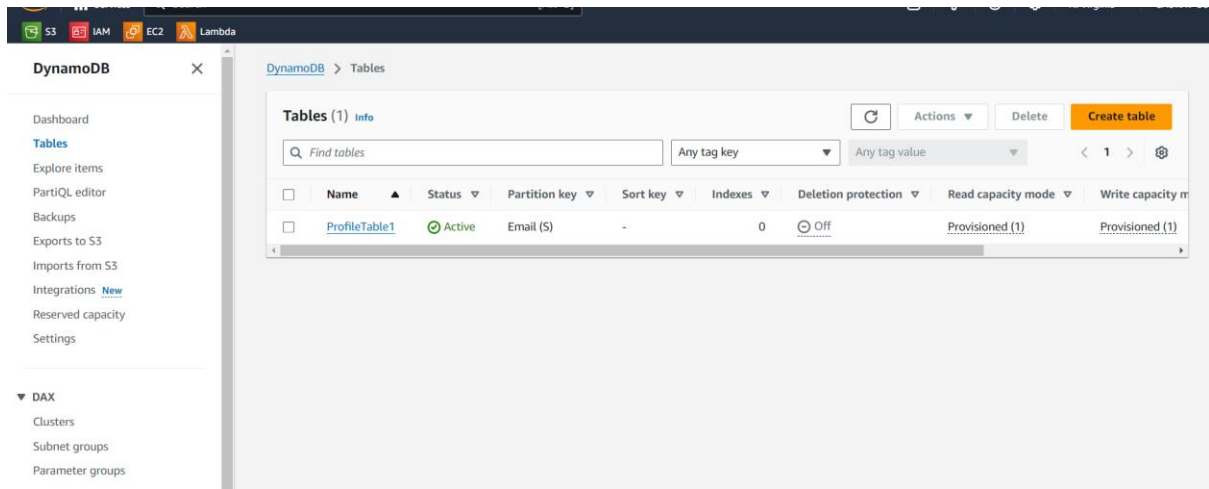
Go to Anything (Ctrl-P)

Environment

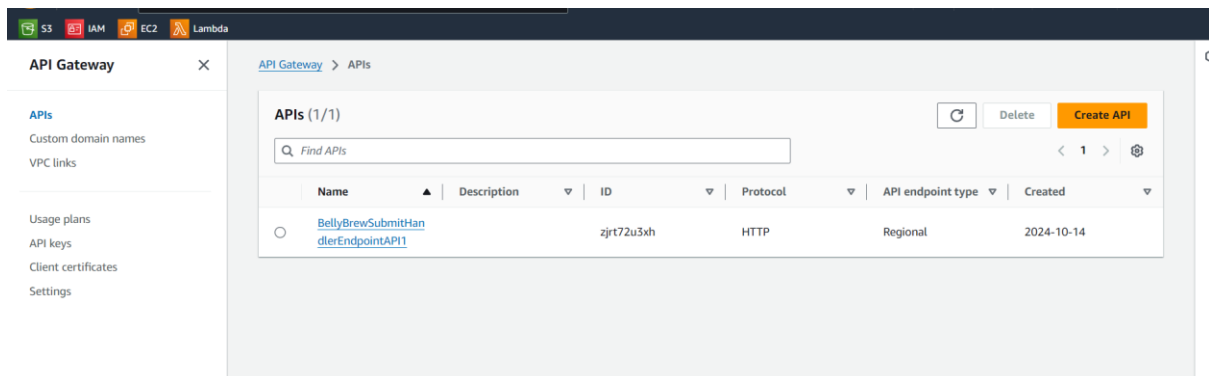
- BellyBrewSubmitHandlerFunction-1
- lambda_function.py

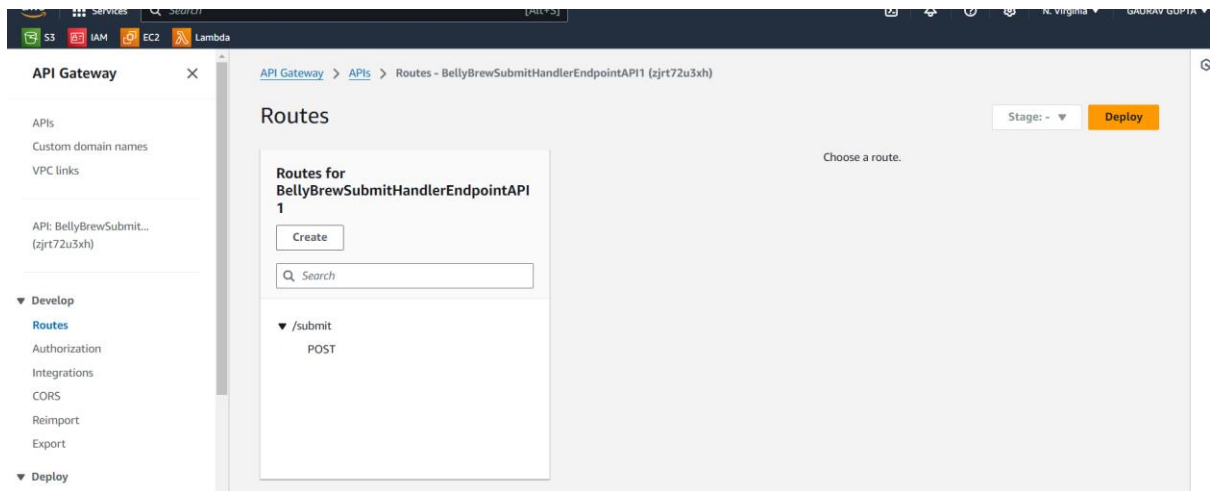
```
1 import os
2 import base64
3 import boto3 # type: ignore
4 from urllib import parse
5 from typing import Any, Dict, Optional
6
7 db: Any = boto3.client("dynamodb") # type: ignore
8 # Explicit region must be set to address a bug in S3/boto3.
9 bucket_region: str = os.environ.get("BOOKSTORE_BUCKET_REGION", "")
10 s3: Any = boto3.client("s3", region_name=bucket_region, endpoint_url=f"https://{bucket_region}.amazonaws.com") # type: ignore
11 ses: Any = boto3.client("sesv2") # type: ignore
12
13
14 def extract_profile_from_event(event):
15     """Extract body which may be in base64 encoding, from API Gateway event."""
16     body_raw = event.get("body", "")
17     is_base64_encoded = event.get("isBase64Encoded", False)
18
19     if is_base64_encoded:
20         body = base64.b64decode(body_raw).decode("ascii")
21     else:
22         body = body_raw
23
24     payload = parse.parse_qs(body, False)
25
26     # The payload dict is formatted as:
27     # {
28     #   'name': ['Jane'],
29     #   'email': ['jane@example.com']
30     # }
31     # The array exists because HTML forms can have multiple <input /> tags with the same "name".
32     # In our case we will only use one <input> per field, so let's use the first element.
```

DynamoDB



API Gateway





SES

